

UNDERSTANDING APACHE AVRO: AVRO INTRODUCTION FOR BIG DATA AND DATA STREAMING ARCHITECTURES

May 8, 2017



Share

Tweet

Like 49

Share

Avro Introduction for Big Data and Data Streaming Architectures

Apache Avro™ is a data serialization system. Avro provides data structures, binary data format, container file format to store persistent data, and provides RPC capabilities. Avro does not require code generation to use and integrates well with JavaScript, Python, Ruby, C, C#, C++ and Java. Avro gets used in the *Hadoop ecosystem* as well as by *Kafka*.

Avro is similar to Thrift, Protocol Buffers, JSON, etc. Avro does not require code generation. Avro needs less encoding as part of the data since it stores names and types in the schema reducing duplication. Avro supports the evolution of schemas.

Why Avro for Kafka and Hadoop?

Avro supports direct mapping to JSON as well as a compact binary format. It is a very fast serialization format. Avro is widely used in the Hadoop ecosystem. Avro supports polyglot bindings to many programming languages and a code generation for static languages. For dynamically typed languages, code generation is not needed. Another key advantage of Avro is its support of evolutionary schemas which supports compatibility checks, and allows evolving your data over time.

Cloudurable provides Kafka training (<http://cloudurable.com/kafka-training/index.html>), Kafka consulting (<http://cloudurable.com/kafka-aws-consulting/index.html>), Kafka support (http://cloudurable.com/subscription_support/index.html) and helps setting up Kafka clusters in AWS (<http://cloudurable.com/services/index.html>).

Avro supports platforms like Kafka that has multiple Producers and Consumers which evolve over time. Avro schemas help keep your data clean and robust.

There was a trend towards schema-less as part of the NoSQL, but that pendulum has swung back a bit e.g., Cassandra has schemas REST/JSON was schema-less and IDL-less but not anymore with Swagger, API gateways, and RAML. Now the trend is more towards schemas that can evolve and Avro fits well in this space.

Avro Schema provides Future Proof Robustness

Streaming architecture like Kafka supports decoupling by sending data in streams to an unknown number of consumers. Streaming architecture is challenging as Consumers and Producers evolve on different timelines. Producers send a stream of records that zero to many Consumers read. Not only are there multiple consumers but data might end up in Hadoop or some other store and used for use cases you did not even imagine. Schemas help future proof your data and make it more robust. Supporting all use cases future (Big Data), past (older Consumers) and current use cases is not easy without a schema. Avro schema with its support for evolution is essential for making the data robust for streaming architectures like Kafka, and with the metadata that schema provides, you can reason on the data. Having a schema provides robustness in providing meta-data about the data stored in Avro records which are self-documenting the data.

Avro provides future usability of data

Data record format compatibility is a hard problem to solve with streaming architecture and Big Data. Avro schemas are not a cure-all, but essential for documenting and modeling your data. Avro Schema definitions capture a point in time of what your data looked like when it recorded since the schema is saved with the data. Data will evolve. New fields are added. Since streams often get recorded in data lakes like Hadoop and those records can represent historical data, not operational data, it makes sense that data streams and data lakes have a less rigid, more evolving schema than the schema of the operational relational database or Cassandra cluster. It makes sense to have a rigid schema for operational data, but not data that ends up in a data lake.

With a streaming platform, consumers and producers can change all of the time and evolve quite a bit. Producers can have Consumers that they never know. You can't test a Consumer that you don't know. For agility sake, you don't want to update every Consumer every time a Producer adds a field to a Record. These types of updates are not feasible without support for Schema.

Avro Schema

Avro data format (wire format and file format) is defined by Avro schemas. When deserializing data, the schema is used. Data is serialized based on the schema, and schema is sent with data or in the case of files stored with the data. Avro data plus schema is fully self-describing data format.

When Avro files store data it also stores schema. Avro RPC is also based on schema, and IDL. Part of the RPC protocol exchanges schemas as part of the handshake. Avro schemas and IDL are written in JSON.

Let's take a look at an example Avro schema.

./src/main/avro/com/cloudurable/phonebook/Employee.avsc

Example schema for an Employee record

```
{ "namespace": "com.cloudurable.phonebook",  
  "type": "record",  
  "name": "Employee",  
  "fields": [  
    { "name": "firstName", "type": "string",  
      { "name": "lastName", "type": "string",  
        { "name": "age", "type": "int",  
          { "name": "phoneNumber", "type": "string",  
            ]  
  }  
}
```

The above defines an employee record with firstName, lastName, age and phoneNumber.

Avro schema generation tools

Avro comes with a set of tools for generating Java classes for Avro types that you define in Avro schema. There are plugins for Maven and Gradle to generate code based on Avro schemas.

This `gradle-avro-plugin` is a Gradle plugin that uses Avro tools to do Java code generation for Apache Avro. This plugin supports Avro schema files (`.avsc`), and Avro RPC IDL (`.avdl`). For Kafka (<http://cloudurable.com/kafka-training/index.html>) you only need `avsc` schema files.

build.gradle - example using gradle-avro-plugin

```
plugins {  
    id "com.commercehub.gradle.plugin.avro" version "0.9.0"  
}  
  
group 'cloudurable'  
version '1.0-SNAPSHOT'  
apply plugin: 'java'  
sourceCompatibility = 1.8  
  
dependencies {  
    compile "org.apache.avro:avro:1.8.1"  
    testCompile group: 'junit', name: 'junit', version: '4.11'  
}  
  
repositories {  
    jcenter()  
    mavenCentral()  
}  
  
avro {  
    createSetters = false  
    fieldVisibility = "PRIVATE"  
}
```

Notice that we did not generate setter methods, and we made the fields private. This makes the instances somewhat immutable.

Running `gradle build` will generate the `Employee.java`.

`./build/generated-main-avro-java/com/cloudurable/phonebook/Employee.java`

Generated Avro code

```
package com.cloudurable.phonebook;

import org.apache.avro.specific.SpecificData;

@SuppressWarnings("all")
@org.apache.avro.specific.AvroGenerated
public class Employee extends org.apache.avro.specific.SpecificRecordBase implements org.apache.avro.specific.SpecificRecord {
    private static final long serialVersionUID = -6112285611684054927L;
    public static final org.apache.avro.Schema SCHEMA$ = new
        org.apache.avro.Schema.Parser().parse("{\"type\":\"record\",\"name\":\"Employee\"..."});
    public static org.apache.avro.Schema getClassSchema() { return SCHEMA$; }
    private java.lang.String firstName;
    private java.lang.String lastName;
    private int age;
    private java.lang.String phoneNumber;
    ...
}
```

The gradle plugin calls the Avro utilities which generates the files and puts them under `build/generated-main-avro-java`

Let's use the generated class as follows to construct an Employee instance.

Using the new Employee class

```
Employee bob = Employee.newBuilder().setAge(35)
    .setFirstName("Bob")
    .setLastName("Jones")
    .setPhoneNumber("555-555-1212")
    .build();

assertEquals("Bob", bob.getFirstName());
```

The Employee class has a constructor and has a builder. We can use the builder to build a new Employee instance.

Next we want to write the Employees to disk.

Writing a list of employees to an Avro file

```
final List<Employee> employeeList = ...
final DatumWriter<Employee> datumWriter = new SpecificDatumWriter<>(Employee.class);
final DataFileWriter<Employee> dataFileWriter = new DataFileWriter<>(datumWriter);

try {
    dataFileWriter.create(employeeList.get(0).getSchema(),
        new File("employees.avro"));
    employeeList.forEach(employee -> {
        try {
            dataFileWriter.append(employee);
        } catch (IOException e) {
            throw new RuntimeException(e);
        }
    });
} finally {
    dataFileWriter.close();
}
```

The above shows serializing an Employee list to disk. In Kafka, we will not be writing to disk directly. We are just showing how so you have a way to test Avro serialization, which is helpful when debugging schema incompatibilities. Note we create a `DatumWriter`, which converts Java instance into an in-memory serialized format. `SpecificDatumWriter` is used with generated classes like `Employee`. `DataFileWriter` writes the serialized records to the `employee.avro` file.

Now let's demonstrate how to read data from an Avro file.

Reading a list of employees from an avro file

```
final File file = new File("employees.avro");
final List<Employee> employeeList = new ArrayList<>();
final DatumReader<Employee> empReader = new SpecificDatumReader<>(Employee.class);
final DataFileReader<Employee> dataFileReader = new DataFileReader<>(file, empReader);

while (dataFileReader.hasNext()) {
    employeeList.add(dataFileReader.next(new Employee()));
}
```

The above deserializes employees from the `employees.avro` file into a `java.util.List` of `Employee` instances. Deserializing is similar to serializing but in reverse. We create a `SpecificDatumReader` to convert in-memory serialized items into instances of our generated `Employee` class. The `DatumReader` reads records from the file by calling `next`. Another way to read is using `forEach` as follows:

Reading a list of employees from an avro file using `forEach`

```
final DataFileReader<Employee> dataFileReader = new DataFileReader<>(file, empReader);  
dataFileReader.forEach(employeeList::add);
```

You can use a `GenericRecord` instead of generating an `Employee` class as follows.

Using `GenericRecord` to create an `Employee` record

```
final String schemaLoc = "src/main/avro/com/cloudurable/phonebook/Employee.avsc";  
final File schemaFile = new File(schemaLoc);  
final Schema schema = new Schema.Parser().parse(schemaFile);  
  
GenericRecord bob = new GenericData.Record(schema);  
bob.put("firstName", "Bob");  
bob.put("lastName", "Smith");  
bob.put("age", 35);  
assertEquals("Bob", bob.get("firstName"));
```

You can write to Avro files using `GenericRecords` as well.

Writing `GenericRecords` to an Avro file

```
final List<GenericRecord> employeeList = new ArrayList<>();

final DatumWriter<GenericRecord> datumWriter = new GenericDatumWriter<>(schema);
final DataFileWriter<GenericRecord> dataFileWriter = new DataFileWriter<>(datumWriter);

try {
    dataFileWriter.create(employeeList.get(0).getSchema(),
        new File("employees2.avro"));
    employeeList.forEach(employee -> {
        try {
            dataFileWriter.append(employee);
        } catch (IOException e) {
            throw new RuntimeException(e);
        }
    });
} finally {
    dataFileWriter.close();
}
```

You can read from Avro files using `GenericRecord`s as well.

Reading GenericRecords from an Avro file

```
final File file = new File("employees2.avro");
final List<GenericRecord> employeeList = new ArrayList<>();
final DatumReader<GenericRecord> empReader = new GenericDatumReader<>();
final DataFileReader<GenericRecord> dataFileReader = new DataFileReader<>(file, empReader);

while (dataFileReader.hasNext()) {
    employeeList.add(dataFileReader.next(null));
}

employeeList.forEach(System.out::println);
```

Avro will validate the data types when it serializes and deserializes the data.

Using the wrong type


```
GenericRecord employee = new GenericData.Record(schema);
employee.put("firstName", "Bob" + index);
employee.put("lastName", "Smith"+ index);
//employee.put("age", index % 35 + 25);
employee.put("age", "OLD");
```

Stack trace from above

```
org.apache.avro.file.DataFileWriter$AppendWriteException: java.lang.ClassCastException:
java.lang.String cannot be cast to java.lang.Number

    at org.apache.avro.file.DataFileWriter.append(DataFileWriter.java:308)
    at com.cloudurable.phonebook.EmployeeTestNoGen.lambda$testWrite$1(EmployeeTestNoGen.java:71)
    at java.util.ArrayList.forEach(ArrayList.java:1249)
    at com.cloudurable.phonebook.EmployeeTestNoGen.testWrite(EmployeeTestNoGen.java:69)
    ...
Caused by: java.lang.ClassCastException: java.lang.String cannot be cast to java.lang.Number
    at org.apache.avro.generic.GenericDatumWriter.writeWithoutConversion(GenericDatumWriter.java:117)
    at org.apache.avro.generic.GenericDatumWriter.write(GenericDatumWriter.java:73)
    at org.apache.avro.generic.GenericDatumWriter.writeField(GenericDatumWriter.java:153)
    at org.apache.avro.generic.GenericDatumWriter.writeRecord(GenericDatumWriter.java:143)
    at org.apache.avro.generic.GenericDatumWriter.writeWithoutConversion(GenericDatumWriter.java:105)
    at org.apache.avro.generic.GenericDatumWriter.write(GenericDatumWriter.java:73)
    at org.apache.avro.generic.GenericDatumWriter.write(GenericDatumWriter.java:60)
    at org.apache.avro.file.DataFileWriter.append(DataFileWriter.java:302)
```

If you left out a required field like `firstName`, then you would get this.

Stack trace from leaving out firstName

```
Caused by: java.lang.NullPointerException: null of string in field firstName of com.cloudurable.phonebook.Employee
    at org.apache.avro.generic.GenericDatumWriter.npe(GenericDatumWriter.java:132)
    at org.apache.avro.generic.GenericDatumWriter.writeWithoutConversion(GenericDatumWriter.java:126)
    at org.apache.avro.generic.GenericDatumWriter.write(GenericDatumWriter.java:73)
    at org.apache.avro.generic.GenericDatumWriter.write(GenericDatumWriter.java:60)
```

In the Avro schema, you can define Records, Arrays, Enums, Unions, Maps and you can use primitive types like String, Int, Boolean, Decimal, Timestamp, Date, and more.

The Avro schema and IDL specification document (<https://avro.apache.org/docs/current/spec.html#Protocol+Declaration>) describes all of the supported types.

Let's add to the Employee schema and show some of the different types that Avro supports.

####

```
{ "namespace": "com.cloudurable.phonebook",
  "type": "record",
  "name": "Employee",
  "fields": [
    { "name": "firstName", "type": "string" },
    { "name": "nickName", "type": [ "null", "string" ], "default" : null },
    { "name": "lastName", "type": "string" },
    { "name": "age", "type": "int" },
    { "name": "emails", "default": [], "type": { "type": "array", "items": "string" } },
    { "name": "phoneNumber", "type":
      [ "null",
        { "type": "record", "name": "PhoneNumber",
          "fields": [
            { "name": "areaCode", "type": "string" },
            { "name": "countryCode", "type": "string", "default" : "" },
            { "name": "prefix", "type": "string" },
            { "name": "number", "type": "string" }
          ]
        }
      ]
    },
    { "name": "status", "default" : "SALARY", "type": { "type": "enum", "name": "Status",
      "symbols" : [ "RETIRED", "SALARY", "HOURLY", "PART_TIME" ] }
    }
  ]
}
```

Avro record attributes are as follows:

- name: name of the record (required).
- namespace: equates to packages or modules

- doc: documentation for future user of this schema
- aliases: array aliases (alias names)
- fields: an array of fields

Avro field attributes are as follows:

- name: name of the field (required)
- doc: description of field (important for future usage)
- type: JSON object defining a schema, or a JSON string naming a record definition (required)
- default: Default value for this field
- order: specifies sort ordering of record (optional, ascending, descending, ignore)
 - aliases: array of alternate names

The `doc` attribute is imperative for future usage as it documents what the fields and records are supposed to represent. Remember that this data can outlive systems that produced it. A self-documenting schema is critical for a robust system.

The above has examples of default values, arrays, primitive types, Records within records, enums, and more.

PhoneNumber record

```
package com.cloudurable.phonebook;

import org.apache.avro.specific.SpecificData;

@SuppressWarnings("all")
@org.apache.avro.specific.AvroGenerated
public class PhoneNumber extends org.apache.avro.specific.SpecificRecordBase ...{
    private static final long serialVersionUID = -3138777939618426199L;
    public static final org.apache.avro.Schema SCHEMA$ =
        new org.apache.avro.Schema.Parser().parse("{\"type\":\"record\",\"name\":\"...
    public static org.apache.avro.Schema getClassSchema() { return SCHEMA$; }
    private java.lang.String areaCode;
    private java.lang.String countryCode;
    private java.lang.String prefix;
    private java.lang.String number;
```

Status enum

```
package com.cloudurable.phonebook;  
@SuppressWarnings("all")  
@org.apache.avro.specific.AvroGenerated  
public enum Status {  
    RETIRED, SALARY, HOURLY, PART_TIME ;  
    ...  
}
```

Tips for using Avro with Kafka and Hadoop

Avoid advanced Avro features which are not supported by polyglot language mappings. Think simple data transfer objects or structs. Don't use magic strings, use enums instead as they provide better validation.

Document all records and fields in the schema. Documentation is imperative for future usage. Documents what the fields and records represent. A self-documenting schema is critical for a robust streaming system and Big Data. Don't use complex union types. Use Unions for nullable fields only and avoid using recursive types at all costs.

Use reasonable field names and use them consistently with other records. Example, `employee_id` instead of `id` and then use `employee_id` in all other records that have a field that refer to the `employee_id` from Employee.

Conclusion

Avro provides fast, compact data serialization. It supports data structures like Records, Maps, Array, and basic types. You can use it direct or use Code Generation. Avro allows schema support to Kafka which we will demonstrate in another article.

Enjoy this slide deck about Avro (<http://cloudurable.com/ppt/avro.pdf>).



[Cassandra / Kafka Support in EC2/AWS](#), [Kafka Training](#), [Kafka Consulting](#)



Avro

Avro

Apache Avro Data
Serialization

1 of 17

Avro Tutorial - Records with Schema for Kafka and Hadoop (<http://www.slideshare.net/JeanPaulAzar1/avro-tutorial-records-with-schema-for-kafka-and-hadoop>) from **Jean-Paul Azar** (<https://www.slideshare.net/JeanPaulAzar1>)

Related content

- What is Kafka? (<http://cloudurable.com/blog/what-is-kafka/index.html>)
- Kafka Architecture (<http://cloudurable.com/blog/kafka-architecture/index.html>)
- Kafka Topic Architecture (<http://cloudurable.com/blog/kafka-architecture-topics/index.html>)
- Kafka Consumer Architecture (<http://cloudurable.com/blog/kafka-architecture-consumers/index.html>)
- Kafka Producer Architecture (<http://cloudurable.com/blog/kafka-architecture-producers/index.html>)
- Kafka Architecture and low level design (<http://cloudurable.com/blog/kafka-architecture-low-level/index.html>)

- [Kafka and Schema Registry \(http://cloudurable.com/blog/kafka-avro-schema-registry/index.html\)](http://cloudurable.com/blog/kafka-avro-schema-registry/index.html)
- [Kafka and Avro \(http://cloudurable.com/blog/avro/index.html\)](http://cloudurable.com/blog/avro/index.html)
- [Kafka Ecosystem \(http://cloudurable.com/blog/kafka-ecosystem/index.html\)](http://cloudurable.com/blog/kafka-ecosystem/index.html)
- [Kafka vs. JMS \(http://cloudurable.com/blog/kafka-vs-jms/index.html\)](http://cloudurable.com/blog/kafka-vs-jms/index.html)
- [Kafka versus Kinesis \(http://cloudurable.com/blog/kinesis-vs-kafka/index.html\)](http://cloudurable.com/blog/kinesis-vs-kafka/index.html)
- [Kafka Tutorial: Using Kafka from the command line \(http://cloudurable.com/blog/kafka-tutorial-kafka-from-command-line/index.html\)](http://cloudurable.com/blog/kafka-tutorial-kafka-from-command-line/index.html)
- [Kafka Tutorial: Kafka Broker Failover and Consumer Failover \(http://cloudurable.com/blog/kafka-tutorial-kafka-failover-kafka-cluster/index.html\)](http://cloudurable.com/blog/kafka-tutorial-kafka-failover-kafka-cluster/index.html)
- [Kafka Tutorial \(http://cloudurable.com/ppt/kafka-tutorial-cloudurable-v2.pdf\)](http://cloudurable.com/ppt/kafka-tutorial-cloudurable-v2.pdf)
- [Kafka Tutorial: Writing a Kafka Producer example in Java \(http://cloudurable.com/blog/kafka-tutorial-kafka-producer/index.html\)](http://cloudurable.com/blog/kafka-tutorial-kafka-producer/index.html)
- [Kafka Tutorial: Writing a Kafka Consumer example in Java \(http://cloudurable.com/blog/kafka-tutorial-kafka-consumer/index.html\)](http://cloudurable.com/blog/kafka-tutorial-kafka-consumer/index.html)
- [Kafka Architecture: Log Compaction \(http://cloudurable.com/blog/kafka-architecture-log-compaction/index.html\)](http://cloudurable.com/blog/kafka-architecture-log-compaction/index.html)
- [Kafka Architecture: Low-Level PDF Slides \(http://cloudurable.com/ppt/4-kafka-detailed-architecture.pdf\)](http://cloudurable.com/ppt/4-kafka-detailed-architecture.pdf)

About Cloudurable

We hope you enjoyed this article. Please provide feedback (<http://cloudurable.com/contact/index.html>). Cloudurable provides Kafka training (<http://cloudurable.com/kafka-training/index.html>), Kafka consulting (<http://cloudurable.com/kafka-aws-consulting/index.html>), Kafka support (http://cloudurable.com/subscription_support/index.html) and helps setting up Kafka clusters in AWS (<http://cloudurable.com/services/index.html>).

Check out our new *GoLang* course. We provide onsite Go Lang training which is instructor led (<http://cloudurable.com/golang-onsite-instructor-led-training/index.html>).



Share

Tweet

Like 49

Share

SEARCH

Search



SHARE

Tweet



Share

facebook

49

Like

Share

FOLLOW

Follow @cloudurable



Follow

64

facebook

CATEGORIES

amazon-ebs (1) (<http://cloudurable.com/categories/amazon-ebs/index.html>)

amazon-ec2 (1) (<http://cloudurable.com/categories/amazon-ec2/index.html>)

amazon-vpc (1) (<http://cloudurable.com/categories/amazon-vpc/index.html>)

ansible (4) (<http://cloudurable.com/categories/ansible/index.html>)

avro (2) (<http://cloudurable.com/categories/avro/index.html>)

aws (4) (<http://cloudurable.com/categories/aws/index.html>)

aws-cassandra (6) (<http://cloudurable.com/categories/aws-cassandra/index.html>)

aws-command-line (1) (<http://cloudurable.com/categories/aws-command-line/index.html>)

cassandra (12) (<http://cloudurable.com/categories/cassandra/index.html>)

cassandra-aws (3) (<http://cloudurable.com/categories/cassandra-aws/index.html>)

cassandra-cluster (1) (<http://cloudurable.com/categories/cassandra-cluster/index.html>)

cassandra-database (2) (<http://cloudurable.com/categories/cassandra-database/index.html>)

cassandra-training (5) (<http://cloudurable.com/categories/cassandra-training/index.html>)

cassandra-tutorial (5) (<http://cloudurable.com/categories/cassandra-tutorial/index.html>)

cloud (4) (<http://cloudurable.com/categories/cloud/index.html>)

cloudformation (1) (<http://cloudurable.com/categories/cloudformation/index.html>)

cloudurable (15) (<http://cloudurable.com/categories/cloudurable/index.html>)

cluster (1) (<http://cloudurable.com/categories/cluster/index.html>)

devops (16) (<http://cloudurable.com/categories/devops/index.html>)

ebs (3) (<http://cloudurable.com/categories/ebs/index.html>)

ec2 (1) (<http://cloudurable.com/categories/ec2/index.html>)

kafka (13) (<http://cloudurable.com/categories/kafka/index.html>)

kafka-advanced-consumers (1) (<http://cloudurable.com/categories/kafka-advanced-consumers/index.html>)

kafka-architecture (13) (<http://cloudurable.com/categories/kafka-architecture/index.html>)

kafka-avro-serialization (1) (<http://cloudurable.com/categories/kafka-avro-serialization/index.html>)

kafka-consulting (2) (<http://cloudurable.com/categories/kafka-consulting/index.html>)

kafka-consumer (1) (<http://cloudurable.com/categories/kafka-consumer/index.html>)

kafka-consumers (1) (<http://cloudurable.com/categories/kafka-consumers/index.html>)

kafka-ecosystem (1) (<http://cloudurable.com/categories/kafka-ecosystem/index.html>)

kafka-schema-registry (1) (<http://cloudurable.com/categories/kafka-schema-registry/index.html>)

kafka-training (23) (<http://cloudurable.com/categories/kafka-training/index.html>)

kafka-tutorial (20) (<http://cloudurable.com/categories/kafka-tutorial/index.html>)

kaka-replication (1) (<http://cloudurable.com/categories/kaka-replication/index.html>)

kinesis (1) (<http://cloudurable.com/categories/kinesis/index.html>)

kinesis-consulting (1) (<http://cloudurable.com/categories/kinesis-consulting/index.html>)

linux (1) (<http://cloudurable.com/categories/linux/index.html>)

metricsd (1) (<http://cloudurable.com/categories/metricsd/index.html>)

microservices (3) (<http://cloudurable.com/categories/microservices/index.html>)

nodetool (1) (<http://cloudurable.com/categories/nodetool/index.html>)

schema-registry (1) (<http://cloudurable.com/categories/schema-registry/index.html>)

smack (3) (<http://cloudurable.com/categories/smack/index.html>)

spark (3) (<http://cloudurable.com/categories/spark/index.html>)

spark--cassandra (1) (<http://cloudurable.com/categories/spark--cassandra/index.html>)

spark--kafka (1) (<http://cloudurable.com/categories/spark--kafka/index.html>)

spark-training (3) (<http://cloudurable.com/categories/spark-training/index.html>)

spark-tutorial (3) (<http://cloudurable.com/categories/spark-tutorial/index.html>)

ssh (1) (<http://cloudurable.com/categories/ssh/index.html>)

ssh-config (1) (<http://cloudurable.com/categories/ssh-config/index.html>)

ssl (1) (<http://cloudurable.com/categories/ssl/index.html>)

systemd (1) (<http://cloudurable.com/categories/systemd/index.html>)

tls (1) (<http://cloudurable.com/categories/tls/index.html>)

vagrant (5) (<http://cloudurable.com/categories/vagrant/index.html>)

TAGS

AKKA ([HTTP://CLOUDURABLE.COM/TAGS/AKKA/INDEX.HTML](http://cloudurable.com/tags/akka/index.html))

AKKA-CONSULTING ([HTTP://CLOUDURABLE.COM/TAGS/AKKA-CONSULTING/INDEX.HTML](http://cloudurable.com/tags/akka-consulting/index.html))

AMAZON-EBS ([HTTP://CLOUDURABLE.COM/TAGS/AMAZON-EBS/INDEX.HTML](http://cloudurable.com/tags/amazon-ebs/index.html))

AMAZON-EC2 ([HTTP://CLOUDURABLE.COM/TAGS/AMAZON-EC2/INDEX.HTML](http://cloudurable.com/tags/amazon-ec2/index.html))

AMI ([HTTP://CLOUDURABLE.COM/TAGS/AMI/INDEX.HTML](http://cloudurable.com/tags/ami/index.html))

ANSIBLE ([HTTP://CLOUDURABLE.COM/TAGS/ANSIBLE/INDEX.HTML](http://cloudurable.com/tags/ansible/index.html))

AVRO ([HTTP://CLOUDURABLE.COM/TAGS/AVRO/INDEX.HTML](http://cloudurable.com/tags/avro/index.html))

AVRO-KAFKA ([HTTP://CLOUDURABLE.COM/TAGS/AVRO-KAFKA/INDEX.HTML](http://cloudurable.com/tags/avro-kafka/index.html))

AWS ([HTTP://CLOUDURABLE.COM/TAGS/AWS/INDEX.HTML](http://cloudurable.com/tags/aws/index.html))

- 🔗 [AWS-CASSANDRA \(HTTP://CLOUDURABLE.COM/TAGS/AWS-CASSANDRA/INDEX.HTML\)](http://cloudurable.com/tags/aws-cassandra/index.html)
- 🔗 [AWS-COMMAND-LINE \(HTTP://CLOUDURABLE.COM/TAGS/AWS-COMMAND-LINE/INDEX.HTML\)](http://cloudurable.com/tags/aws-command-line/index.html)
- 🔗 [CASSANDRA \(HTTP://CLOUDURABLE.COM/TAGS/CASSANDRA/INDEX.HTML\)](http://cloudurable.com/tags/cassandra/index.html)
- 🔗 [CASSANDRA-ARCHITECTURE \(HTTP://CLOUDURABLE.COM/TAGS/CASSANDRA-ARCHITECTURE/INDEX.HTML\)](http://cloudurable.com/tags/cassandra-architecture/index.html)
- 🔗 [CASSANDRA-AWS \(HTTP://CLOUDURABLE.COM/TAGS/CASSANDRA-AWS/INDEX.HTML\)](http://cloudurable.com/tags/cassandra-aws/index.html)
- 🔗 [CASSANDRA-CLOUD \(HTTP://CLOUDURABLE.COM/TAGS/CASSANDRA-CLOUD/INDEX.HTML\)](http://cloudurable.com/tags/cassandra-cloud/index.html)
- 🔗 [CASSANDRA-CLUSTER \(HTTP://CLOUDURABLE.COM/TAGS/CASSANDRA-CLUSTER/INDEX.HTML\)](http://cloudurable.com/tags/cassandra-cluster/index.html)
- 🔗 [CASSANDRA-DATABASE \(HTTP://CLOUDURABLE.COM/TAGS/CASSANDRA-DATABASE/INDEX.HTML\)](http://cloudurable.com/tags/cassandra-database/index.html)
- 🔗 [CASSANDRA-DBA \(HTTP://CLOUDURABLE.COM/TAGS/CASSANDRA-DBA/INDEX.HTML\)](http://cloudurable.com/tags/cassandra-dba/index.html)
- 🔗 [CASSANDRA-DEVOPS \(HTTP://CLOUDURABLE.COM/TAGS/CASSANDRA-DEVOPS/INDEX.HTML\)](http://cloudurable.com/tags/cassandra-devops/index.html)
- 🔗 [CASSANDRA-OS-SYSTEM-MEMORY \(HTTP://CLOUDURABLE.COM/TAGS/CASSANDRA-OS-SYSTEM-MEMORY/INDEX.HTML\)](http://cloudurable.com/tags/cassandra-os-system-memory/index.html)
- 🔗 [CASSANDRA-TRAINING \(HTTP://CLOUDURABLE.COM/TAGS/CASSANDRA-TRAINING/INDEX.HTML\)](http://cloudurable.com/tags/cassandra-training/index.html)
- 🔗 [CASSANDRA-TUTORIAL \(HTTP://CLOUDURABLE.COM/TAGS/CASSANDRA-TUTORIAL/INDEX.HTML\)](http://cloudurable.com/tags/cassandra-tutorial/index.html)
- 🔗 [CLOUD \(HTTP://CLOUDURABLE.COM/TAGS/CLOUD/INDEX.HTML\)](http://cloudurable.com/tags/cloud/index.html)
- 🔗 [CLOUDFORMATION \(HTTP://CLOUDURABLE.COM/TAGS/CLOUDFORMATION/INDEX.HTML\)](http://cloudurable.com/tags/cloudformation/index.html)
- 🔗 [CLOUDFORMATION-TUTORIAL \(HTTP://CLOUDURABLE.COM/TAGS/CLOUDFORMATION-TUTORIAL/INDEX.HTML\)](http://cloudurable.com/tags/cloudformation-tutorial/index.html)
- 🔗 [CLOUDURABLE \(HTTP://CLOUDURABLE.COM/TAGS/CLOUDURABLE/INDEX.HTML\)](http://cloudurable.com/tags/cloudurable/index.html)
- 🔗 [CLUSTER \(HTTP://CLOUDURABLE.COM/TAGS/CLUSTER/INDEX.HTML\)](http://cloudurable.com/tags/cluster/index.html)

🔗 [COMPUTE \(HTTP://CLOUDURABLE.COM/TAGS/COMPUTE/INDEX.HTML\)](http://cloudurable.com/tags/compute/index.html)

🔗 [CONSUMERS \(HTTP://CLOUDURABLE.COM/TAGS/CONSUMERS/INDEX.HTML\)](http://cloudurable.com/tags/consumers/index.html)

🔗 [DBA \(HTTP://CLOUDURABLE.COM/TAGS/DBA/INDEX.HTML\)](http://cloudurable.com/tags/dba/index.html)

🔗 [DEVOPS \(HTTP://CLOUDURABLE.COM/TAGS/DEVOPS/INDEX.HTML\)](http://cloudurable.com/tags/devops/index.html)

🔗 [EBS \(HTTP://CLOUDURABLE.COM/TAGS/EBS/INDEX.HTML\)](http://cloudurable.com/tags/ebs/index.html)

🔗 [EC2 \(HTTP://CLOUDURABLE.COM/TAGS/EC2/INDEX.HTML\)](http://cloudurable.com/tags/ec2/index.html)

🔗 [EC2-INSTANCE-STORE \(HTTP://CLOUDURABLE.COM/TAGS/EC2-INSTANCE-STORE/INDEX.HTML\)](http://cloudurable.com/tags/ec2-instance-store/index.html)

🔗 [ECU \(HTTP://CLOUDURABLE.COM/TAGS/ECU/INDEX.HTML\)](http://cloudurable.com/tags/ecu/index.html)

🔗 [FAILOVER \(HTTP://CLOUDURABLE.COM/TAGS/FAILOVER/INDEX.HTML\)](http://cloudurable.com/tags/failover/index.html)

🔗 [IMAGES \(HTTP://CLOUDURABLE.COM/TAGS/IMAGES/INDEX.HTML\)](http://cloudurable.com/tags/images/index.html)

🔗 [INSTANCES \(HTTP://CLOUDURABLE.COM/TAGS/INSTANCES/INDEX.HTML\)](http://cloudurable.com/tags/instances/index.html)

🔗 [JMS \(HTTP://CLOUDURABLE.COM/TAGS/JMS/INDEX.HTML\)](http://cloudurable.com/tags/jms/index.html)

🔗 [KAFKA \(HTTP://CLOUDURABLE.COM/TAGS/KAFKA/INDEX.HTML\)](http://cloudurable.com/tags/kafka/index.html)

🔗 [KAFKA-ADVANCED-CONSUMERS \(HTTP://CLOUDURABLE.COM/TAGS/KAFKA-ADVANCED-CONSUMERS/INDEX.HTML\)](http://cloudurable.com/tags/kafka-advanced-consumers/index.html)

🔗 [KAFKA-ADVANCED-PRODUCERS \(HTTP://CLOUDURABLE.COM/TAGS/KAFKA-ADVANCED-PRODUCERS/INDEX.HTML\)](http://cloudurable.com/tags/kafka-advanced-producers/index.html)

🔗 [KAFKA-ARCHITECTURE \(HTTP://CLOUDURABLE.COM/TAGS/KAFKA-ARCHITECTURE/INDEX.HTML\)](http://cloudurable.com/tags/kafka-architecture/index.html)

🔗 [KAFKA-AVRO \(HTTP://CLOUDURABLE.COM/TAGS/KAFKA-AVRO/INDEX.HTML\)](http://cloudurable.com/tags/kafka-avro/index.html)

🔗 [KAFKA-CONNECT \(HTTP://CLOUDURABLE.COM/TAGS/KAFKA-CONNECT/INDEX.HTML\)](http://cloudurable.com/tags/kafka-connect/index.html)

🔗 [KAFKA-CONSULTING \(HTTP://CLOUDURABLE.COM/TAGS/KAFKA-CONSULTING/INDEX.HTML\)](http://cloudurable.com/tags/kafka-consulting/index.html)

🔗 [KAFKA-CONSUMERS \(HTTP://CLOUDURABLE.COM/TAGS/KAFKA-CONSUMERS/INDEX.HTML\)](http://cloudurable.com/tags/kafka-consumers/index.html)

🔗 [KAFKA-CONSUMERS-ADVANCED \(HTTP://CLOUDURABLE.COM/TAGS/KAFKA-CONSUMERS-ADVANCED/INDEX.HTML\)](http://cloudurable.com/tags/kafka-consumers-advanced/index.html)

🔗 [KAFKA-ECOSYSTEM \(HTTP://CLOUDURABLE.COM/TAGS/KAFKA-ECOSYSTEM/INDEX.HTML\)](http://cloudurable.com/tags/kafka-ecosystem/index.html)

🔗 [KAFKA-LOG-COMPACTION \(HTTP://CLOUDURABLE.COM/TAGS/KAFKA-LOG-COMPACTION/INDEX.HTML\)](http://cloudurable.com/tags/kafka-log-compaction/index.html)

🔗 [KAFKA-REST-PROXY \(HTTP://CLOUDURABLE.COM/TAGS/KAFKA-REST-PROXY/INDEX.HTML\)](http://cloudurable.com/tags/kafka-rest-proxy/index.html)

🔗 [KAFKA-STREAMS \(HTTP://CLOUDURABLE.COM/TAGS/KAFKA-STREAMS/INDEX.HTML\)](http://cloudurable.com/tags/kafka-streams/index.html)

🔗 [KAFKA-TRAINING \(HTTP://CLOUDURABLE.COM/TAGS/KAFKA-TRAINING/INDEX.HTML\)](http://cloudurable.com/tags/kafka-training/index.html)

🔗 [KAFKA-TUTORIAL \(HTTP://CLOUDURABLE.COM/TAGS/KAFKA-TUTORIAL/INDEX.HTML\)](http://cloudurable.com/tags/kafka-tutorial/index.html)

🔗 [KINESIS \(HTTP://CLOUDURABLE.COM/TAGS/KINESIS/INDEX.HTML\)](http://cloudurable.com/tags/kinesis/index.html)

🔗 [KINESIS-CONSULTING \(HTTP://CLOUDURABLE.COM/TAGS/KINESIS-CONSULTING/INDEX.HTML\)](http://cloudurable.com/tags/kinesis-consulting/index.html)

🔗 [LINUX \(HTTP://CLOUDURABLE.COM/TAGS/LINUX/INDEX.HTML\)](http://cloudurable.com/tags/linux/index.html)

🔗 [METRICSD \(HTTP://CLOUDURABLE.COM/TAGS/METRICSD/INDEX.HTML\)](http://cloudurable.com/tags/metricSD/index.html)

🔗 [MICROSERVICES \(HTTP://CLOUDURABLE.COM/TAGS/MICROSERVICES/INDEX.HTML\)](http://cloudurable.com/tags/microservices/index.html)

🔗 [MICROSERVICES-ARCHITECTURE \(HTTP://CLOUDURABLE.COM/TAGS/MICROSERVICES-ARCHITECTURE/INDEX.HTML\)](http://cloudurable.com/tags/microservices-architecture/index.html)

🔗 [NAT \(HTTP://CLOUDURABLE.COM/TAGS/NAT/INDEX.HTML\)](http://cloudurable.com/tags/nat/index.html)

🔗 [NODETOOL \(HTTP://CLOUDURABLE.COM/TAGS/NODETOOL/INDEX.HTML\)](http://cloudurable.com/tags/nodeTool/index.html)

🔗 [NUMA \(HTTP://CLOUDURABLE.COM/TAGS/NUMA/INDEX.HTML\)](http://cloudurable.com/tags/numa/index.html)

🔗 [PRODUCERS \(HTTP://CLOUDURABLE.COM/TAGS/PRODUCERS/INDEX.HTML\)](http://cloudurable.com/tags/producers/index.html)

🔗 [QBIT \(HTTP://CLOUDURABLE.COM/TAGS/QBIT/INDEX.HTML\)](http://cloudurable.com/tags/qbit/index.html)

🔗 [RAM \(HTTP://CLOUDURABLE.COM/TAGS/RAM/INDEX.HTML\)](http://cloudurable.com/tags/ram/index.html)

🔗 [REAKT \(HTTP://CLOUDURABLE.COM/TAGS/REAKT/INDEX.HTML\)](http://cloudurable.com/tags/reakt/index.html)

🔗 [REPLICATION \(HTTP://CLOUDURABLE.COM/TAGS/REPLICATION/INDEX.HTML\)](http://cloudurable.com/tags/replication/index.html)

🔗 [SCHEMA-REGISTRY \(HTTP://CLOUDURABLE.COM/TAGS/SCHEMA-REGISTRY/INDEX.HTML\)](http://cloudurable.com/tags/schema-registry/index.html)

🔗 [SMACK \(HTTP://CLOUDURABLE.COM/TAGS/SMACK/INDEX.HTML\)](http://cloudurable.com/tags/smack/index.html)

🔗 [SPARK \(HTTP://CLOUDURABLE.COM/TAGS/SPARK/INDEX.HTML\)](http://cloudurable.com/tags/spark/index.html)

🔗 **SPARK--CASSANDRA** ([HTTP://CLOUDURABLE.COM/TAGS/SPARK--CASSANDRA/INDEX.HTML](http://cloudurable.com/tags/spark--cassandra/index.html))

🔗 **SPARK-TRAINING** ([HTTP://CLOUDURABLE.COM/TAGS/SPARK-TRAINING/INDEX.HTML](http://cloudurable.com/tags/spark-training/index.html))

🔗 **SPARK-TUTORIAL** ([HTTP://CLOUDURABLE.COM/TAGS/SPARK-TUTORIAL/INDEX.HTML](http://cloudurable.com/tags/spark-tutorial/index.html))

🔗 **SSH** ([HTTP://CLOUDURABLE.COM/TAGS/SSH/INDEX.HTML](http://cloudurable.com/tags/ssh/index.html))

🔗 **SSH-CONFIG** ([HTTP://CLOUDURABLE.COM/TAGS/SSH-CONFIG/INDEX.HTML](http://cloudurable.com/tags/ssh-config/index.html))

🔗 **SSL** ([HTTP://CLOUDURABLE.COM/TAGS/SSL/INDEX.HTML](http://cloudurable.com/tags/ssl/index.html))

🔗 **SYSTEMD** ([HTTP://CLOUDURABLE.COM/TAGS/SYSTEMD/INDEX.HTML](http://cloudurable.com/tags/systemd/index.html))

🔗 **TLS** ([HTTP://CLOUDURABLE.COM/TAGS/TLS/INDEX.HTML](http://cloudurable.com/tags/tls/index.html))

🔗 **VAGRANT** ([HTTP://CLOUDURABLE.COM/TAGS/VAGRANT/INDEX.HTML](http://cloudurable.com/tags/vagrant/index.html))

🔗 **VCPU** ([HTTP://CLOUDURABLE.COM/TAGS/VCPU/INDEX.HTML](http://cloudurable.com/tags/vcpu/index.html))

🔗 **VPC** ([HTTP://CLOUDURABLE.COM/TAGS/VPC/INDEX.HTML](http://cloudurable.com/tags/vpc/index.html))

🔗 **WHAT-IS-KAFKA** ([HTTP://CLOUDURABLE.COM/TAGS/WHAT-IS-KAFKA/INDEX.HTML](http://cloudurable.com/tags/what-is-kafka/index.html))

Apache Spark Training (<http://cloudurable.com/spark-training/index.html>)

Kafka Tutorial (<http://cloudurable.com/blog/kafka-tutorial/index.html>)

Akka Consulting (<http://cloudurable.com/akka-consulting/index.html>)

Cassandra Training (<http://cloudurable.com/cassandra-course/index.html>)

AWS Cassandra Database Support (http://cloudurable.com/subscription_support_benefits_cassandra/index.html)

Kafka Support Pricing (http://cloudurable.com/subscription_support/index.html?q=kafka)

Cassandra Database Support Pricing (http://cloudurable.com/subscription_support/index.html?q=cassandra)

Non-stop Cassandra (<http://cloudurable.com/cloudurable-cassandra-watchdog/index.html?q=cassandra>)

Watchdog (<http://cloudurable.com/cloudurable-cassandra-watchdog/index.html?q=watchdog>)

Advantages of using Cloudurable™ (<http://cloudurable.com/advantages/index.html>)

Cassandra Consulting (<http://cloudurable.com/service-quick-start-mentoring-cassandra-or-kafka-aws-ec2/index.html>)

Cloudurable™| Guide to AWS Cassandra Deploy (<http://cloudurable.com/ppt/amazon-cassandra.pdf>)

Cloudurable™| AWS Cassandra Guidelines and Notes (<http://cloudurable.com/ppt/amazon-cassandra-notes.pdf>)

Free guide to deploying Cassandra on AWS (<http://cloudurable.com/cassandra-aws-consulting/index.html>)

Kafka Training (<http://cloudurable.com/kafka-training/index.html>)

Kafka Consulting (<http://cloudurable.com/kafka-aws-consulting/index.html>)

DynamoDB Training (<http://cloudurable.com/dynamodb-training/index.html>)

DynamoDB Consulting (<http://cloudurable.com/dynamodb-consulting/index.html>)

Kinesis Training (<http://cloudurable.com/kinesis-training/index.html>)

Kinesis Consulting (<http://cloudurable.com/kinesis-consulting/index.html>)

Kafka Tutorial PDF (<http://cloudurable.com/blog/kafka-tutorial-v1/index.html>)

Redis Consulting (<http://cloudurable.com/redis-consulting/index.html>)

Redis Training (<http://cloudurable.com/redis-onsite-instructor-led-training/index.html>)

ElasticSearch / ELK Consulting (<http://cloudurable.com/elk-consulting/index.html>)

ElasticSearch Training (<http://cloudurable.com/elasticsearch-onsite-instructor-led-training/index.html>)

InfluxDB/TICK Training (<http://cloudurable.com/influxdb-onsite-instructor-led-training/index.html>) TICK Consulting (<http://cloudurable.com/tick-consulting/index.html>)

ABOUT US

Cloudurable™: Leader in AWS cloud computing for Kafka™, Cassandra™ Database, Apache Spark, AWS CloudFormation™ DevOps. We do **Cassandra training, Apache Spark, Kafka training, Kafka consulting** and **cassandra consulting** with a focus on AWS and data engineering. (FAQ (<http://cloudurable.com/faq/index.html>))

FOLLOW CLOUDURABLE™

facebook page (<https://www.facebook.com/cloudurable>)

google plus (<https://plus.google.com/116648719730180908239>)

twitter (<https://twitter.com/cloudurable>)

linkedin (<https://www.linkedin.com/company/17964258/>)

Why Cloudurable™?

Advantage of using Cloudurable™ (<http://cloudurable.com/advantages/index.html>)

About Cloudurable™?

About Cloudurable™ (<http://cloudurable.com/faq/index.html>)


What are the benefits of using subscription support?

Benefits of Subscription Cassandra Support (http://cloudurable.com/subscription_support_benefits_cassandra/index.html)

RECENT POSTS

KAFKA CONSUMER: ADVANCED CONSUMERS ([HTTP://CLOUDURABLE.COM/BLOG/KAFKA-ADVANCED-CONSUMER-1/INDEX.HTML](http://cloudurable.com/blog/kafka-advanced-consumer-1/index.html))

KAFKA TUTORIAL ([HTTP://CLOUDURABLE.COM/BLOG/KAFKA-TUTORIAL/INDEX.HTML](http://cloudurable.com/blog/kafka-tutorial/index.html))



KAFKA TUTORIAL: CREATING ADVANCED KAFKA PRODUCERS IN JAVA ([HTTP://CLOUDURABLE.COM/BLOG/KAFKA-TUTORIAL-KAFKA-PRODUCER-ADVANCED-JAVA-EXAMPLES/INDEX.HTML](http://cloudurable.com/blog/kafka-tutorial-kafka-producer-advanced-java-examples/index.html))

KAFKA CONSULTING ([HTTP://CLOUDURABLE.COM/BLOG/KAFKA-CONSULTING/INDEX.HTML](http://cloudurable.com/blog/kafka-consulting/index.html))
[n/blog/kafka-](http://cloudurable.com/blog/kafka-consulting/index.html)
[n/blog/kafka-](http://cloudurable.com/blog/kafka-consulting/index.html)

CONTACT

<http://cloudurable.com/blog/kafka-examples/index.html>
401 California Street

San Francisco

CA 94111

USA

America

(415) 758-1113 (tel:14157581113)

GO TO CONTACT PAGE (/CONTACT/INDEX.HTML)

Copyright © 2015 - 2018, Cloudurable™, all rights reserved. Streamline your Cassandra Database, Apache Spark and Kafka DevOps in AWS. SMACK/Lambda architecture consulting! Spark, Mesos, Akka, Cassandra and Kafka in AWS.

Apache Spark Training (<http://cloudurable.com/spark-training/index.html>), Akka Consulting (<http://cloudurable.com/akka-consulting/index.html>), AWS Cassandra Support (http://cloudurable.com/subscription_support_benefits_cassandra/index.html), Cassandra Training (<http://cloudurable.com/cassandra-course/index.html>), Kafka Training (<http://cloudurable.com/kafka-training/index.html>), Cassandra Consulting (<http://cloudurable.com/service-architecture-analysis-cassandra-or-kafka-aws-ec2/index.html>), Kafka Consulting (<http://cloudurable.com/kafka-aws-consulting/index.html>), Spark Training (<http://cloudurable.com/spark-aws-emr-training/index.html>), Spark Consulting (<http://cloudurable.com/spark-aws-emr-consulting/index.html>), Kafka Tutorial (<http://cloudurable.com/blog/kafka-tutorial-v1/index.html>)

